


Skills Over MCP

From Experiment to Extension

Ola Hungerford · AGNTCon + MCPCon Japan 2026



About me




NORDSTROM

 Model Context Protocol




What an MCP server hands the model

an agentic entity




MCP server



**Tool descriptions say what each tool does,
not how to orchestrate them for a task.**


Order, domain conventions, and conditional logic come from the skill.

There are many ways to build a spaceship



With raw materials and tools alone, the agent explores the solution space.


"Build me a spaceship" might get what you want, or it might not.



tools: a spaceship

tools + skill: the spaceship you asked for


One way to think about Skills over MCP is like shipping the manual in the same box as the tools and other components.




**skills from a repo / plugin,
tools from the MCP server**

**Skills Over MCP: one channel,
one set of access controls**


Solutions the working group discussed, and where we landed




a new primitive
SEP-2076, closed,
kick-started the group




a pointer in server instructions
can work as a complement,
but are statically loaded





bare resources
uri conventions can work,
but no way to list skills



tools that list and read it
for tools-only hosts



on disk




over MCP
listed by skills/list

Skill files as MCP resources, plus a listing.

Several implementations had converged on this before the group chose it.

The model reads context on demand as resources



What each skills manifest entry pins down

docs-server

skill://refunds/SKILL.md

≠

vendor-server

skill://refunds/SKILL.md

Identity is (server, uri)

The name is only a label. Same-named skills from different origins never shadow one another.

SKILL.md

sha256 · size

references/FORMS.md

sha256 · size

templates/invoice.md

sha256 · size


×

references/EXTRA.md

not in the manifest

Every file carries a digest

The set is complete, unless the skill is dynamic. Outside of explicitly dynamic skills, a read that doesn't match, or a file the manifest never listed, is a verification failure.




resources/directory/read


Directory listing, optional

Lists the direct children of a directory, so "pick a template from templates/" works without knowing every URI. Gated behind directoryRead.


The Five Stages of the Skills Extension



The Other Five Stages of the Skills Extension





Key core maintainer feedback




Archives over MCP were too risky, everything needs a digest - **nothing opaque on the wire.**

Adding is easier than removing,
start smaller and try to solve for most of the problem.



Too much optionality fragments the ecosystem,
need to pick an explicit mode for skills ingestion.

Who showed up




Contributors came from 60+ organizations who spoke, wrote, reviewed, or committed

What we haven't solved yet within this extension



Many tools, resources, skills, efficiently?



**Core Primitives WG,
Progressive Disclosure WG**

MCP Roadmap



How do I audit, validate, mutate
skills and other context over MCP?



Interceptors WG

SEP-2624



Files the same remote and local?




Filesystems WG

Chartered



Dependencies and versions?




Agent Skills spec

agentskills.io



Was the skill actually used?




Observability WG

Proposed

Thank you

Come build with us



MCP contributor communication
modelcontextprotocol.io/community/communication